

Wes
McKinney
(pandas)

Francesc Alted
(Señor pytables)

Fernando
Perez
(ipython)

Travis
Oliphant
(numpy,
Continuum)

Paul Ivanov
(matplotlib)



Databases with Python

AY250 Fall 2013



authors: J. Bloom



This work is licensed under a [Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License](#).

Overview of Today's Lecture

- What are DBs? Why DBs?
- Generic DB concepts:
relational table architecture SQL
- sqlite3 & Python
- MySQL and Postgresql
- HDF5 and NetCDF4

What are DBs?

Organized collection of data, managed by a Database Management System (DBMS)

DBMSs (generally) offer:

- storage
- creation & manipulation of data
- search
- (high-level data operations)

Whaa? Doesn't Python already do this for me?

creation
manipulation
search
storage

```
>>> db = range(1,1000,2)
>>> db[10] = "spamalot"
>>> db.index(31)
15
>>> import cPickle ; cPickle.dump(db,open("my.py.db","w"))
```

Limitations:

- Searches on lists are slow, $O(n)$
- db can only be as big as the available RAM
- efficiently storage files (i.e., pickles)
are python-only (portability issue)
- single-instance use only (no concurrent use by others)

dictionaries solve some of simplest the searching issues
(using a hash value for the key), but any realistic search
is still painfully slow on big data:

```
db = {"scenes": \  
      [{"cheeseshop": {"actors": ["Cleese", "Palin"]}}, , \  
       {"spam": {"actors": ["Cleese", "Idle"]}}], , \  
      "actors": \  
      {"Palin": {"age": 54, "how funny": 6}, , \  
       "Cleese": {"age": 62, "how funny": 9}, , \  
       "Idle": {"age": 61, "how funny": None}}}
```

what's the average age of the actors in the cheeseshop scene?

(concept of a JOIN operation...see more later)

Python "Solution" is to 3rd-Party It

provides built-in access (as of 2.5) to sqlite3
& APIs for all other major DBMS

General Purpose Database Systems

- IBM [DB2](#)
- [Firebird](#) (and Interbase)
- [Informix](#)
- [Ingres](#)
- MySQL
- Oracle
- PostgreSQL
- SAP DB (also known as "MaxDB")
- Microsoft [SQL Server](#)
- Sybase

somewhat
outdated:

Record-based Databases

Databases working on flat files or fixed records.

- [MetaKit](#)
- [ZODB](#)
- [BerkeleyDB](#)
- [KirbyBase](#)
- [Durus](#)
- [atop](#)
- [buzhug](#)

XML Databases

- 4Suite server
- Oracle/Sleepycat DB XML

<http://wiki.python.org/moin/DatabaseInterfaces>

Understanding DBMS Frameworks

Three levels

1. External

what "users" of the DB sees (i.e., the Python API)
usually different for different users

2. Conceptual

logical structure of the DB, relationships between the data,
security details, etc. "DB Scheme"

3. Physical

how the data are stored, managed at a low
level (from compiled-language code to bytes)



Implementation

How the conceptual level views data, the "data model":

Relational

Think of data organized as 2D tables,
with connections between them

[most]

Object-Oriented

Think of data as attributes in model
classes, with methods & inheritance

db4o

~BigTable

<http://www.odbms.org/downloads.aspx>

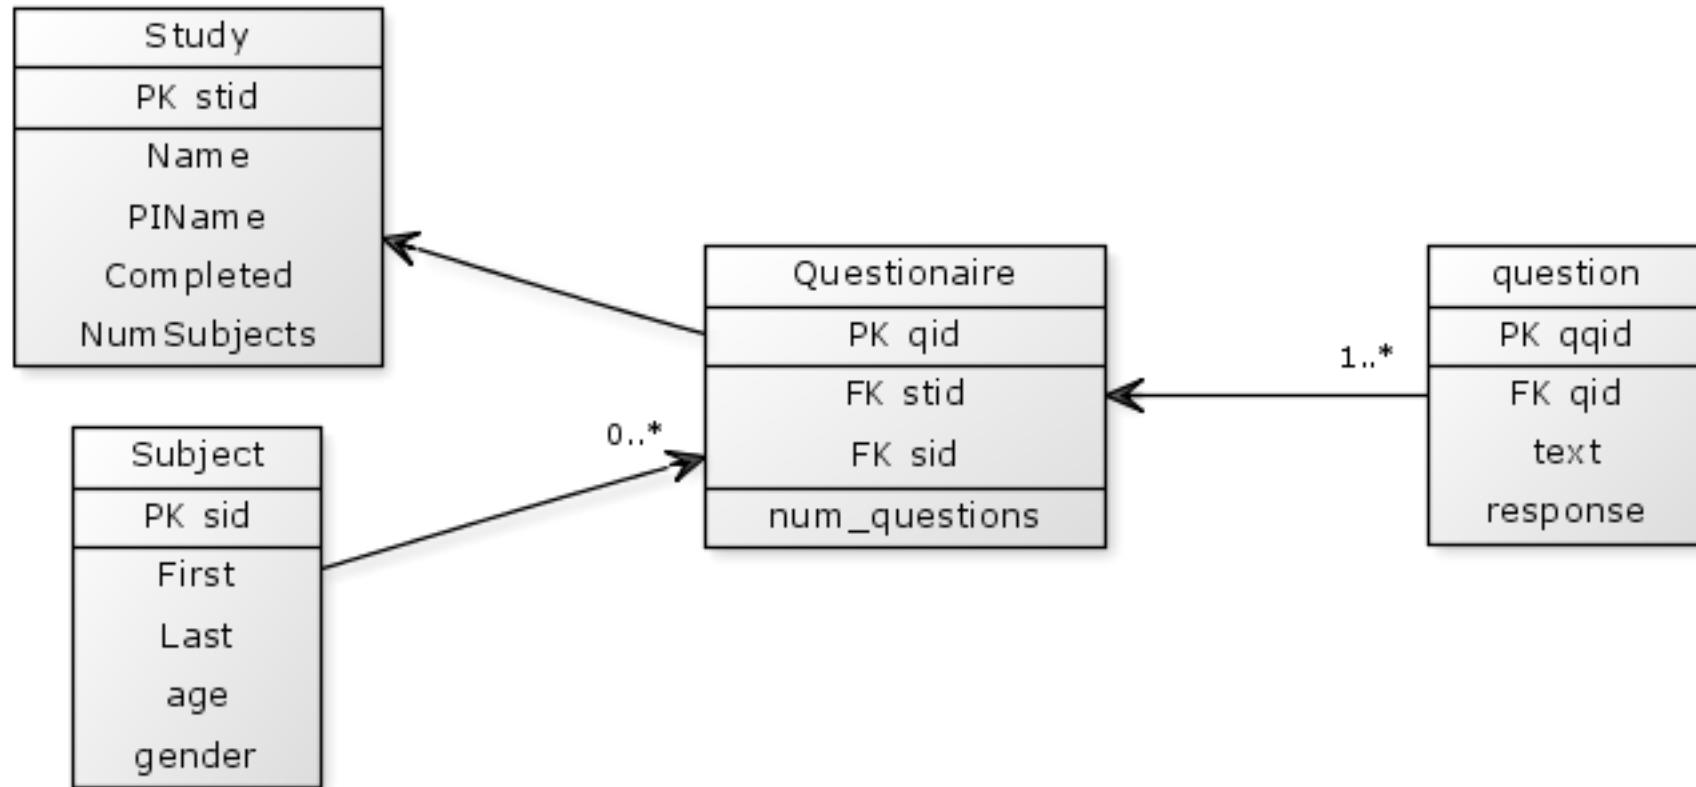
Hierarchical

Tree- or document-like. XML/JSON
DBs and “NoSQL” DBs

eXist, RethinkDB

DynamoDB,
mongoDB
couchDB, ...

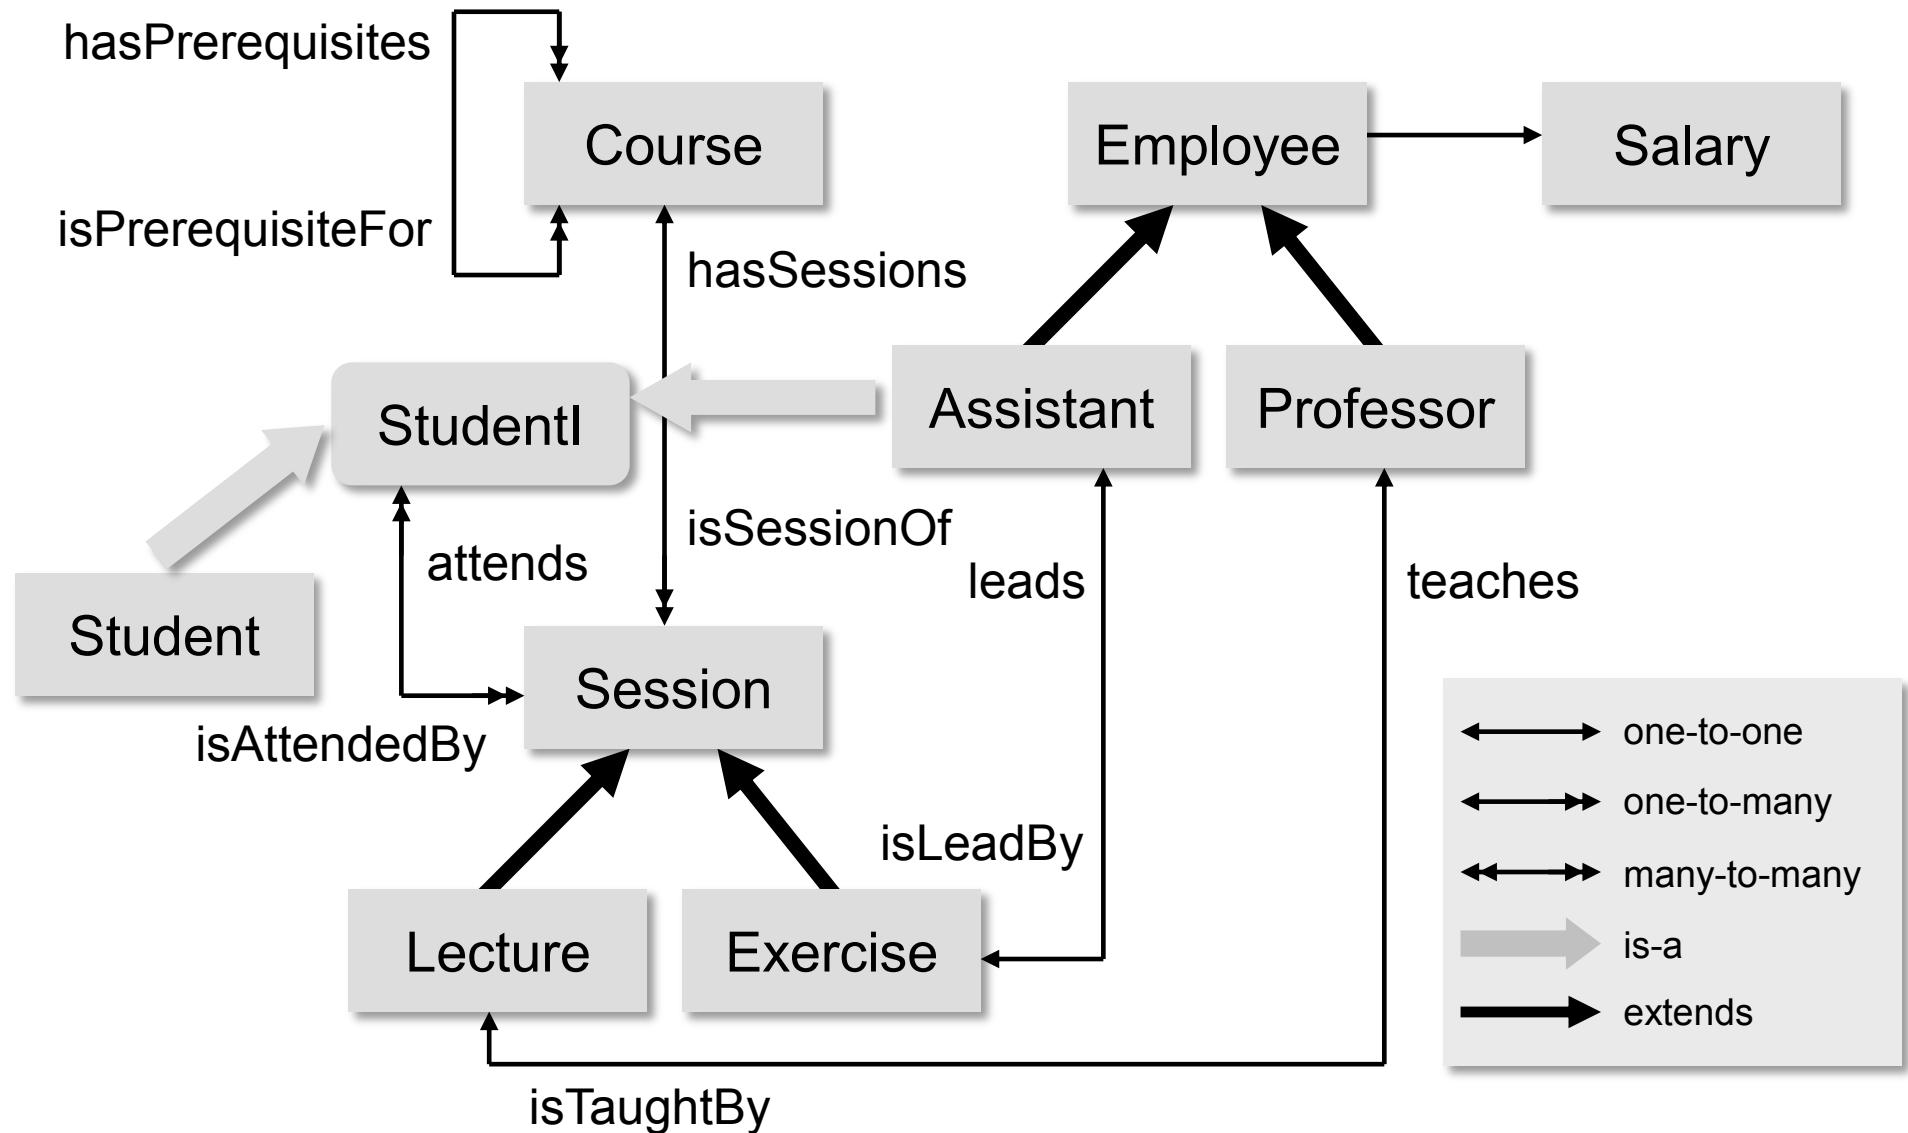
Relational



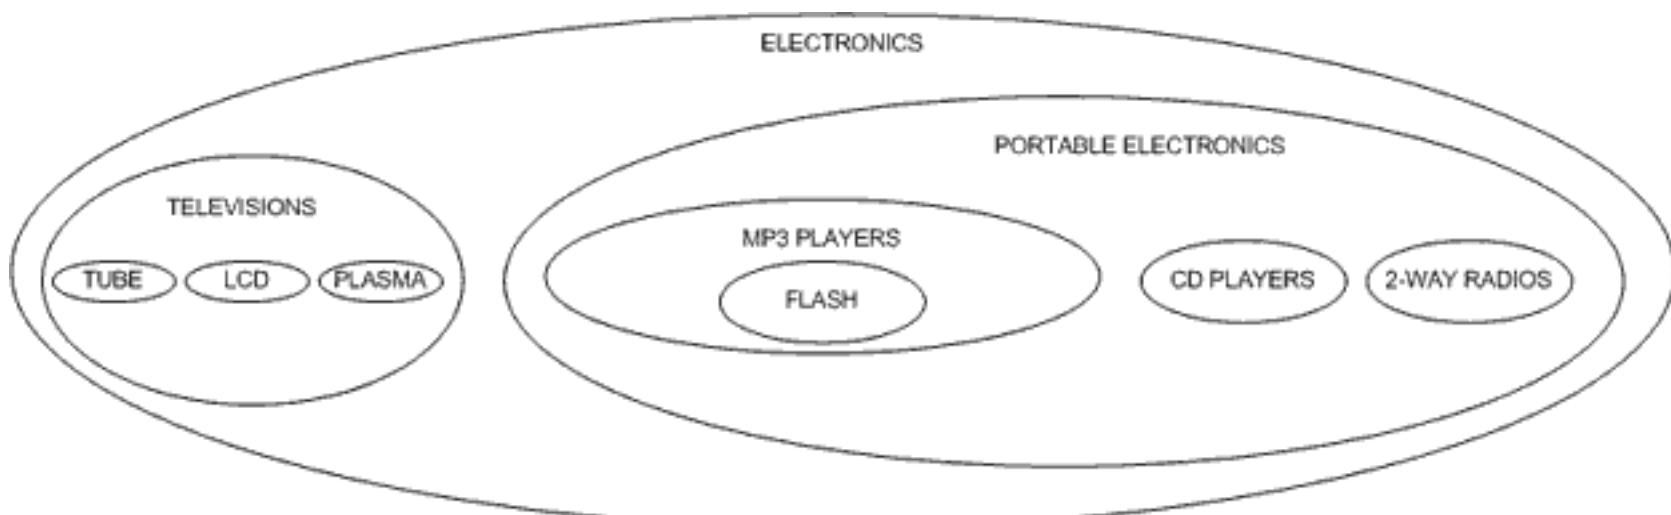
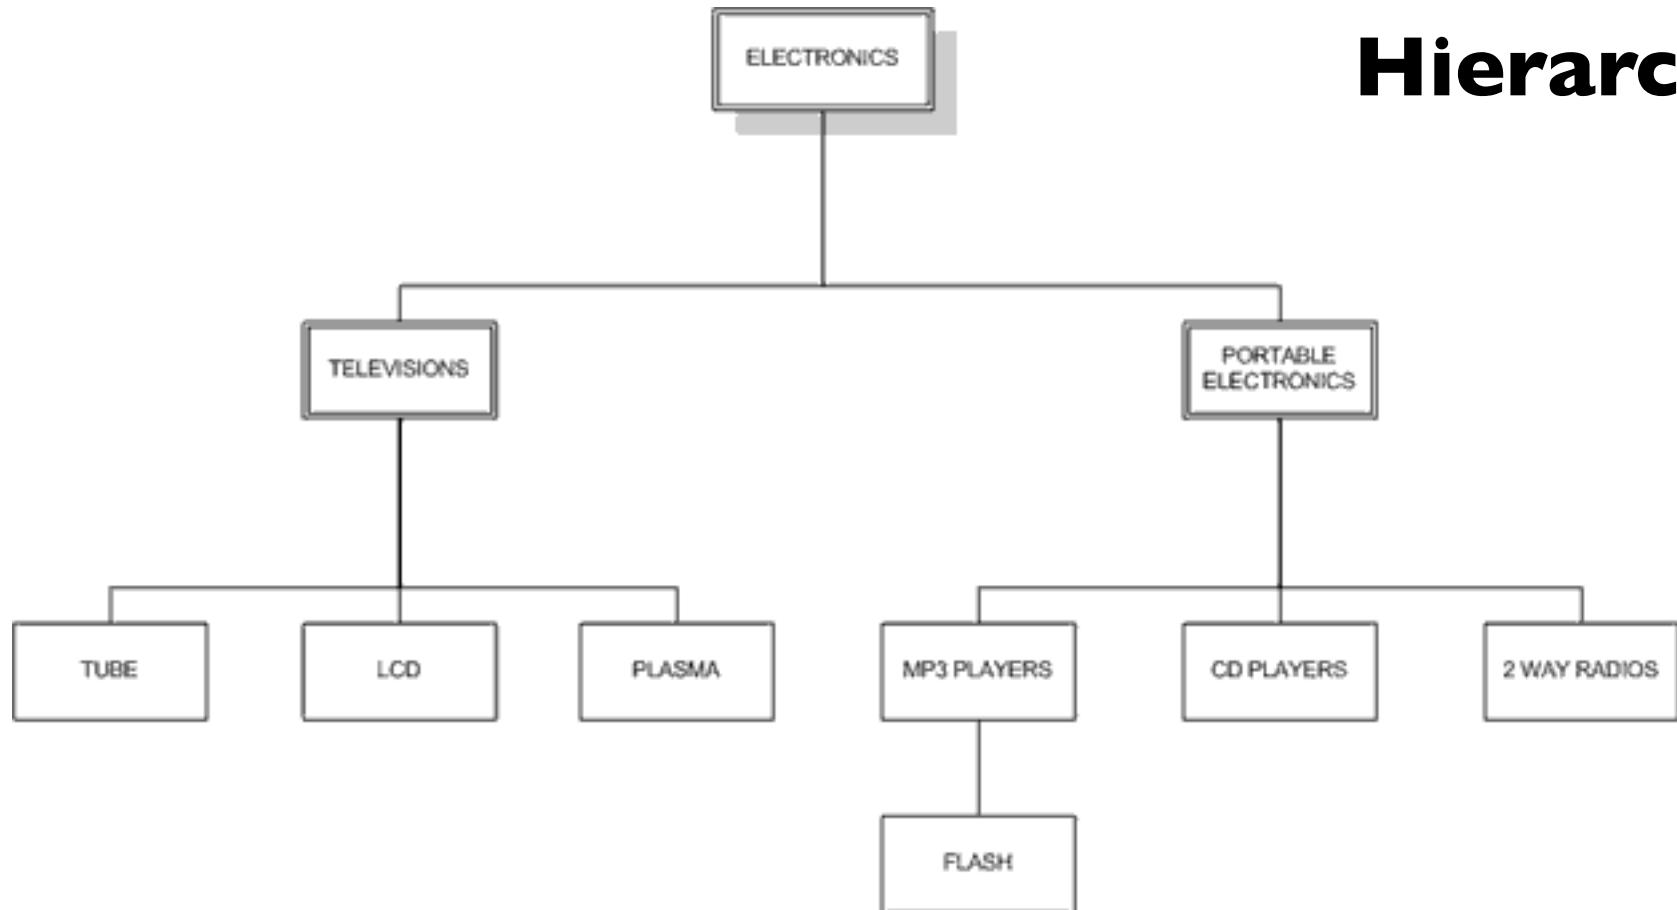
PK: Primary Key (Hashed)

FK: Foreign Key

Object-Oriented



Hierarchical



Implementation

How the DBMS operates:

client-server - persistent server

managing and accepting connections from
multiple clients. (could have a distributed
implementation)

MySQL,
Postgresql,
...
...

embedded - no server. Clients connect to
the DB directly.

sqlite3

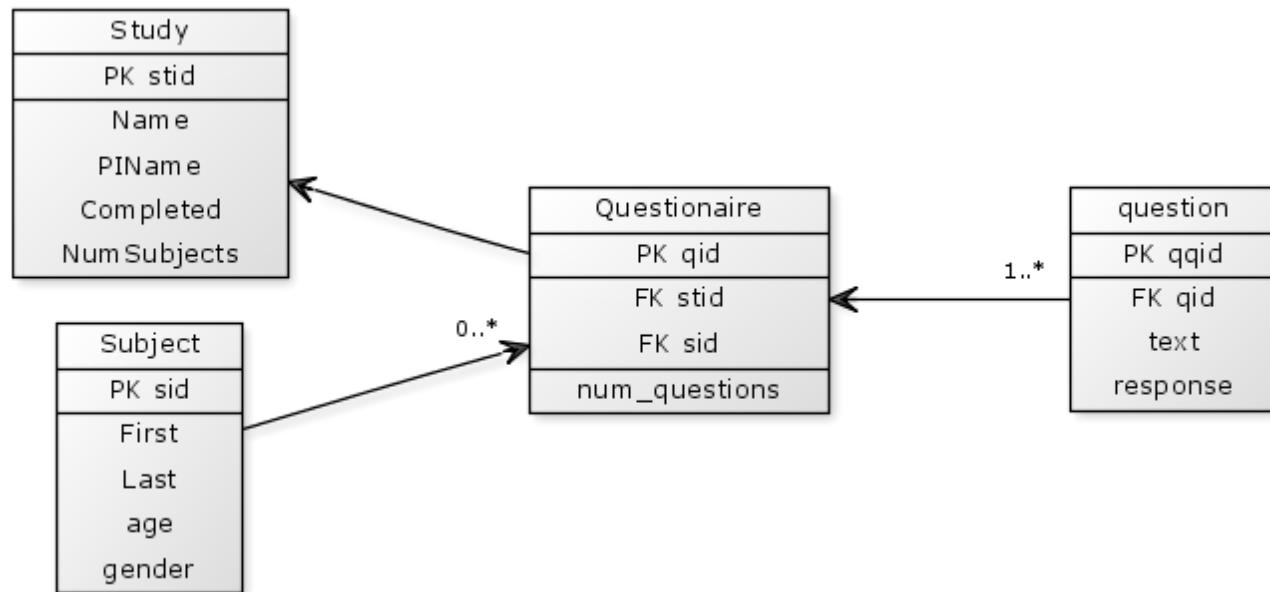
All modern DBMSs are meant to be
fast, safe, very scalable:
large, complex queries can be efficient

Data Model Concepts

database - contains tables, actions ("triggers"), permissions/security

table - data with fixed number of columns and rows of data that you add to (and query)

keys - hashed columns (or combinations of columns) used for efficiently joining of tables. "Primary" & "Foreign"



de Young

FINE ARTS MUSEUMS OF SAN FRANCISCO

Calendar | E-Newsletter | Contact | Blog | Press Room | Log In

search

Exhibitions Visiting Tickets About Collections Programs + Events Education Members Support Museum Store

HOME | COLLECTIONS

Search the Collections

also: <http://www.bampfa.berkeley.edu/search/artcollection>

Search

[Advanced Search](#)

diebenkorn

[reset form](#)

SEARCH

Your Galleries

You are not logged in.

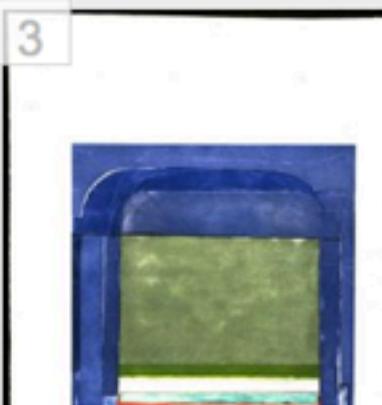
[Log In To My Gallery](#)

or

[Create An Account](#)

1 2 3 4 ... 31 32 33 34

Page 1 of 34. 404 matches found.



Simple (Command-Line) DB Creation

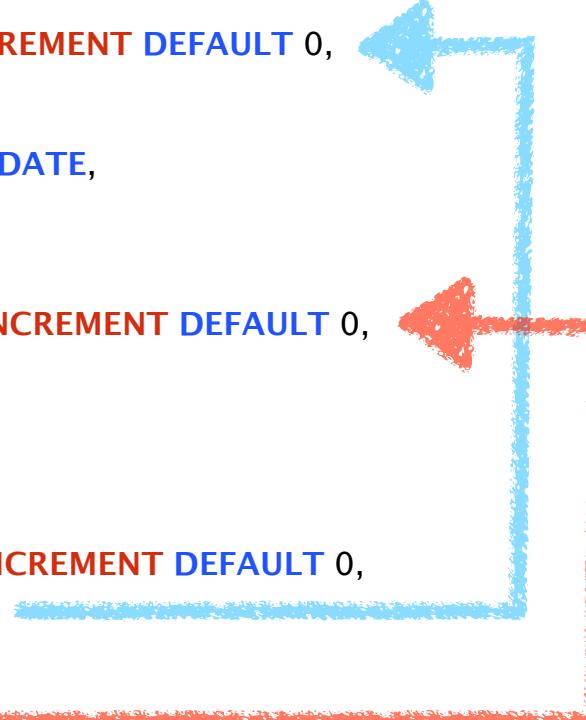
```
PySeminar> sqlite3 art.db
SQLite version 3.7.7.1 2011-06-28 17:39:05
Enter ".help" for instructions
Enter SQL statements terminated with a ";"

sqlite> .read art-create.sql
sqlite> .schema
```

```
CREATE TABLE artist (
    aid integer NOT NULL PRIMARY KEY AUTOINCREMENT DEFAULT 0,
    first_name text,
    last_name text,
    birth_date date NOT NULL DEFAULT CURRENT_DATE,
    birth_country text);

CREATE TABLE museum (
    mid INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT DEFAULT 0,
    name TEXT,
    country TEXT,
    city TEXT);

CREATE TABLE work (
    wid INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT DEFAULT 0,
    aid INTEGER NOT NULL DEFAULT 0,
    title TEXT,
    type TEXT,
    mid integer,
    finish_date Date);
```



"Data Definition Language (DDL)"

Loading in some data

```
PySeminar> sqlite3 art.db
SQLite version 3.7.7.1 2011-06-28 17:39:05
Enter ".help" for instructions
Enter SQL statements terminated with a ";"
sqlite> .read art-load.sql
sqlite> .header on
sqlite> .mode column
sqlite> select * from work;
wid      aid      title      type      mid      finish_date
-----  -----  -----
1        3        Seawall    Painting   2        1957-06-01
2        3        Spreading  Print     1        1981-02-30
3        1        Bedroom at Painting   2        1889-09-03
4        2        Campbell's screenprin 2        1968-09-08
5        2        Liz        lithograph 1        1966-02-22
```

"Data Manipulation Language (DML)":

`select * from tablename;`

Basic Queries in SQL

```
select * from work;
```

gets all rows from table work and prints out all columns

```
select title,type from work;
```

gets all rows from table work and prints selected columns

```
select title,type as "work type" from work;
```

*gets all rows from table work and prints selected columns,
renaming the second one (view only)*

```
select title,DATE('NOW') - finish_date as age  
from work where  
finish_date < DATE("1970-01-01")  
order by age desc;
```

*get only those rows where the work was finished before 1970, and
show the name and age in years, descending order by age*

Basic Queries in SQL

```
sqlite> select * from work where type = "painting";
sqlite>
sqlite> select * from work where type like "painting";
wid      aid      title      type      mid      finish_date
-----  -----  -----
1        3        Seawall    Painting   2        1957-06-01
3        1        Bedroom at Painting   2        1889-09-03
sqlite> select * from work where type glob "P*";
wid      aid      title      type      mid      finish_date
-----  -----  -----
1        3        Seawall    Painting   2        1957-06-01
2        3        Spreading  Print     1        1981-02-30
3        1        Bedroom at Painting   2        1889-09-03
sqlite> select * from work where type glob "*[p|P]rint*";
wid      aid      title      type      mid      finish_date
-----  -----  -----
2        3        Spreading Spade  Print     1        1981-02-30
4        2        Campbell's Toma Screenprin  2        1968-09-08
```

Small exercise:

What work was created less than 35 years ago and starts with the letter "S"?

What if I want to view all works in Berkeley by American artists?

This information is distributed across three different tables

aid	first_name	last_name	birth_date	birth_country	death_date	artist
1	Vincent	Van Gogh	1853-03-30	Netherlands	1890-07-29	
2	Andy	Warhol	1928-08-06	USA	1987-02-22	
3	Richard	Diebenkorn	1922-04-22	USA	1993-03-30	

wid	aid	title	type	mid	finish_date	work
1	3	Seawall	Painting	2	1957-06-01	
2	3	Spreading	Print	1	1981-02-30	
3	1	Bedroom at	Painting	2	1889-09-03	
4	2	Campbell's	screenprin	2	1968-09-08	
5	2	Liz	lithograph	1	1966-02-22	

museum	mid	name	country	city
	1	Berkeley Art	USA	Berkeley
	2	de Young	USA	San Franci

SQL: JOIN (after SELECT, before WHERE)

What if I want to view all works in Berkeley by American artists?

This information is distributed across three different tables

```
sqlite> select museum.name as museum,artist.last_name as artist,work.title,work.type,work.finish_date as date from work  
      JOIN artist on artist.aid = work.aid  
      JOIN museum on museum.mid = work.mid  
      where artist.birth_country = 'USA'  
      and museum.city = 'Berkeley';
```

museum	artist	title	type	date
Berkeley Art	Diebenkorn	Spreading Spade	Print	1981-02-30
Berkeley Art	Warhol	Liz	lithograph	1966-02-22

<http://zetcode.com/databases/sqlitetutorial/>

More Advanced Concepts

Views

"virtual tables" that are the result of a (complex)query on the tables. Like "Smart Folders"

Triggers

built in actions that the DB takes on itself when something is done to it. E.g., log all the insert actions in a separate log table

Transactions

series of DB interactions that can be committed "all at once" or "rolledback" if there is a problem

Practical Database Usage in Python



authors: J. Bloom



Outline

- 1) Built-in Database Module
 - `sqlite3`
 - Plotting from a database
- 2) MySQL, PostgreSQL
 - `MySQLdb`, `PyGreSQL`, others
- 3) Breakout Exercise
 - Plotting seismograph stations
- 4) HDF
 - `tables`, `netCDF4`

sqlite3

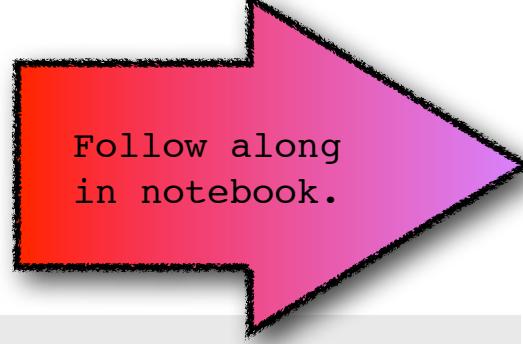
- Built-in SQL database access
- Database is stored as a file (or in RAM)
- Syntax uses Structured Query Language
- Attempts to protect the database from corruption

The next slides will show how to

- create a database
- create a table
- insert some data
- join specific columns of tables
- left join tables

sqlite3

create a database & table



Follow along
in notebook.

```
>>> import sqlite3
>>> connection = sqlite3.connect("/tmp/example.db")
>>> cursor = connection.cursor()
>>> sql_cmd = """CREATE TABLE dan_aykroyd (id INTEGER PRIMARY KEY AUTOINCREMENT,
...     skit_title TEXT, air_date DATE, season INT, ep INT, role TEXT)"""
>>> cursor.execute(sql_cmd)
```

example.db

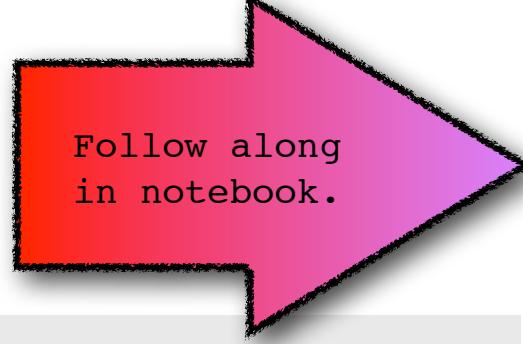
dan_aykroyd

id	skit_title	air_date	season	ep	role

[continue interpreter session next slide]

sqlite3

insert some data



Follow along
in notebook.

```
>>> skit_data = [  
...     ("Trojan Horse Home Security", "10/11/75", 1, 1, "Kenny Vorstrather"),  
...     ("E. Buzz Miller's Animal Kingdom", "2/25/78", 3, 12, "E. Buzz Miller"),  
...     ("The Coneheads at Home", "10/21/78", 4, 3, "Beldar Conehead")]  
>>> for role in skit_data:  
...     sql_cmd = ("INSERT INTO dan_aykroyd (skit_title, air_date, season, " +  
...                 "ep, role) VALUES " + str(role))  
>>> cursor.execute(sql_cmd)
```

example.db

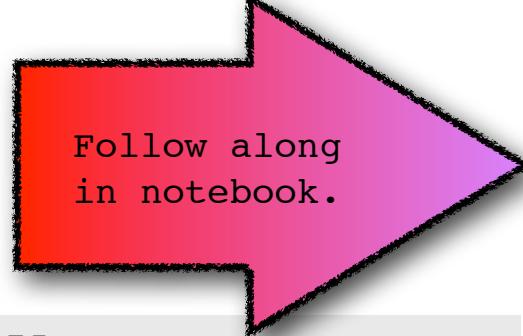
dan_aykroyd

id	skit_title	air_date	season	ep	role
1	"Trojan Horse Home Security"	"10/11/75"	1	1	"Kenny Vorstrather"
2	"E. Buzz Miller's Animal Kingdom"	"2/25/78"	3	12	"E. Buzz Miller"
3	"The Coneheads at Home"	"10/21/78"	4	3	"Beldar Conehead"

[continue interpreter session next slide]

sqlite3

run simple query



Follow along
in notebook.

```
>>> sql_cmd = "SELECT * FROM dan_aykroyd WHERE air_date > '10/12/75'"  
>>> cursor.execute(sql_cmd)  
>>> db_info = cursor.fetchall()  
>>> for entry in db_info: print entry  
    (2, u"E. Buzz Miller's Animal Kingdom", u'2/25/78', 3, 12, u'E. Buzz Miller')  
    (3, u'The Coneheads at Home', u'10/21/78', 4, 3, u'Beldar Conehead')  
>>> connection.commit()  
>>> connection.close()
```

example.db

dan_aykroyd

id	skit_title	air_date	season	ep	role
1	"Trojan Horse Home Security"	"10/11/75"	1	1	"Kenny Vorstrather"
2	"E. Buzz Miller's Animal Kingdom"	"2/25/78"	3	12	"E. Buzz Miller"
3	"The Coneheads at Home"	"10/21/78"	4	3	"Beldar Conehead"

sqlite3

Follow along
in notebook.

create another table, enter data,
run simple join on columns

```
>>> import sqlite3
>>> connection = sqlite3.connect("/tmp/example.db")
>>> cursor = connection.cursor()
>>> sql_cmd = """CREATE TABLE jane_curtin (id INTEGER PRIMARY KEY AUTOINCREMENT,
...     skit_title TEXT, air_date DATE, season INT, episode INT, role TEXT)"""
>>> cursor.execute(sql_cmd)
>>> skit_data = [
...     ("Teen Talk", "7/24/76", 1, 23, "Jane"),
...     ("The Snakehandling O'Sheas", "9/25/76", 2, 2, "Jane O'Shea"),
...     ("The Coneheads at Home", "10/21/78", 4, 3, "Prymaat Conehead")]
>>> for role in skit_data:
...     sql_cmd = ("INSERT INTO jane_curtin (skit_title, air_date, season, " +
...               "episode, role) VALUES " + str(role))
...     cursor.execute(sql_cmd)
>>> sql_cmd = """SELECT dan_aykroyd.skit_title, dan_aykroyd.air_date,
...     dan_aykroyd.season, dan_aykroyd.episode, dan_aykroyd.role, jane_curtin.role
...     FROM dan_aykroyd, jane_curtin WHERE
...     dan_aykroyd.skit_title = jane_curtin.skit_title AND
...     dan_aykroyd.air_date = jane_curtin.air_date"""
>>> cursor.execute(sql_cmd)
>>> db_info = cursor.fetchall()
>>> for entry in db_info: print entry
(u'The Coneheads at Home', u'10/21/78', 4, 3, u'Beldar Conehead', u'Prymaat Conehead')
```

[continue interpreter session next code slide]

sqlite3

join specific columns

Follow along
in notebook.

example.db

dan_aykroyd

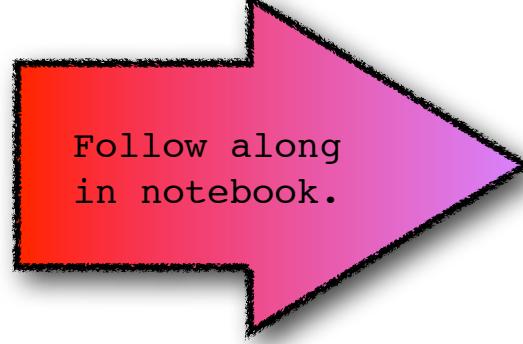
id	skit_title	air_date	season	ep	role
1	"Trojan Horse Home Security"	"10/11/75"	1	1	"Kenny Vorstrather"
2	"E. Buzz Miller's Animal Kingdom"	"2/25/78"	3	12	"E. Buzz Miller"
3	"The Coneheads at Home"	"10/21/78"	4	3	"Beldar Conehead"

jane_curtin

id	skit_title	air_date	season	ep	role
1	"Teen Talk"	"7/24/76"	1	23	"Jane"
2	"The Snakehandling O'Sheas"	"9/25/76"	2	2	"Jane O'Shea"
3	"The Coneheads at Home"	"10/21/78"	4	3	"Prymaat Conehead"

sqlite3

run left join



Follow along
in notebook.

[interpreter session continued from previous code slide]

```
>>> sql_cmd = """SELECT dan_aykroyd.skit_title, dan_aykroyd.air_date,  
...     dan_aykroyd.season, dan_aykroyd.ep, dan_aykroyd.role, jane_curtin.role  
...     FROM dan_aykroyd LEFT JOIN jane_curtin ON  
...     dan_aykroyd.skit_title = jane_curtin.skit_title AND  
...     dan_aykroyd.air_date = jane_curtin.air_date"""  
>>> cursor.execute(sql_cmd)  
>>> db_info = cursor.fetchall()  
>>> for entry in db_info: print entry  
(u'Trojan Horse Home Security', u'10/11/75', 1, 1, u'Kenny Vorstrather', None)  
(u"E. Buzz Miller's Animal Kingdom", u'2/25/78', 3, 12, u'E. Buzz Miller', None)  
(u'The Coneheads at Home', u'10/21/78', 4, 3, u'Beldar Conehead', u'Prymaat Conehead')  
>>> connection.commit()  
>>> connection.close()
```

sqlite3

left join

dan_aykroyd_curtin

id	skit_title	air_date	season	ep	role	role
1	"Teen Talk"	"10/24/75"	1	23	"Kenny Vorstrather"	"Jane"
2	"The Snakehandling O'Sheas"	"2/25/76"	2	12	"E. Buzz Miller"	"Jane O'Shea"
3	"The Coneheads at Home"	"10/21/78"	4	3	"Beldar Conehead"	"Prymaat Conehead"

jane_curtin

id	skit_title	air_date	season	ep	role
1	"Teen Talk"	"7/24/76"	1	23	"Jane"
2	"The Snakehandling O'Sheas"	"9/25/76"	2	2	"Jane O'Shea"
3	"The Coneheads at Home"	"10/21/78"	4	3	"Prymaat Conehead"

left-joined table

id	skit_title	air_date	season	ep	dan_aykroyd.role	jane_curtin.role
1	"Trojan Horse Home Security"	"10/11/75"	1	1	"Kenny Vorstrather"	None
2	"E. Buzz Miller's Animal Kingdom"	"2/25/78"	3	12	"E. Buzz Miller"	None
3	"The Coneheads at Home"	"10/21/78"	4	3	"Beldar Conehead"	"Prymaat Conehead"

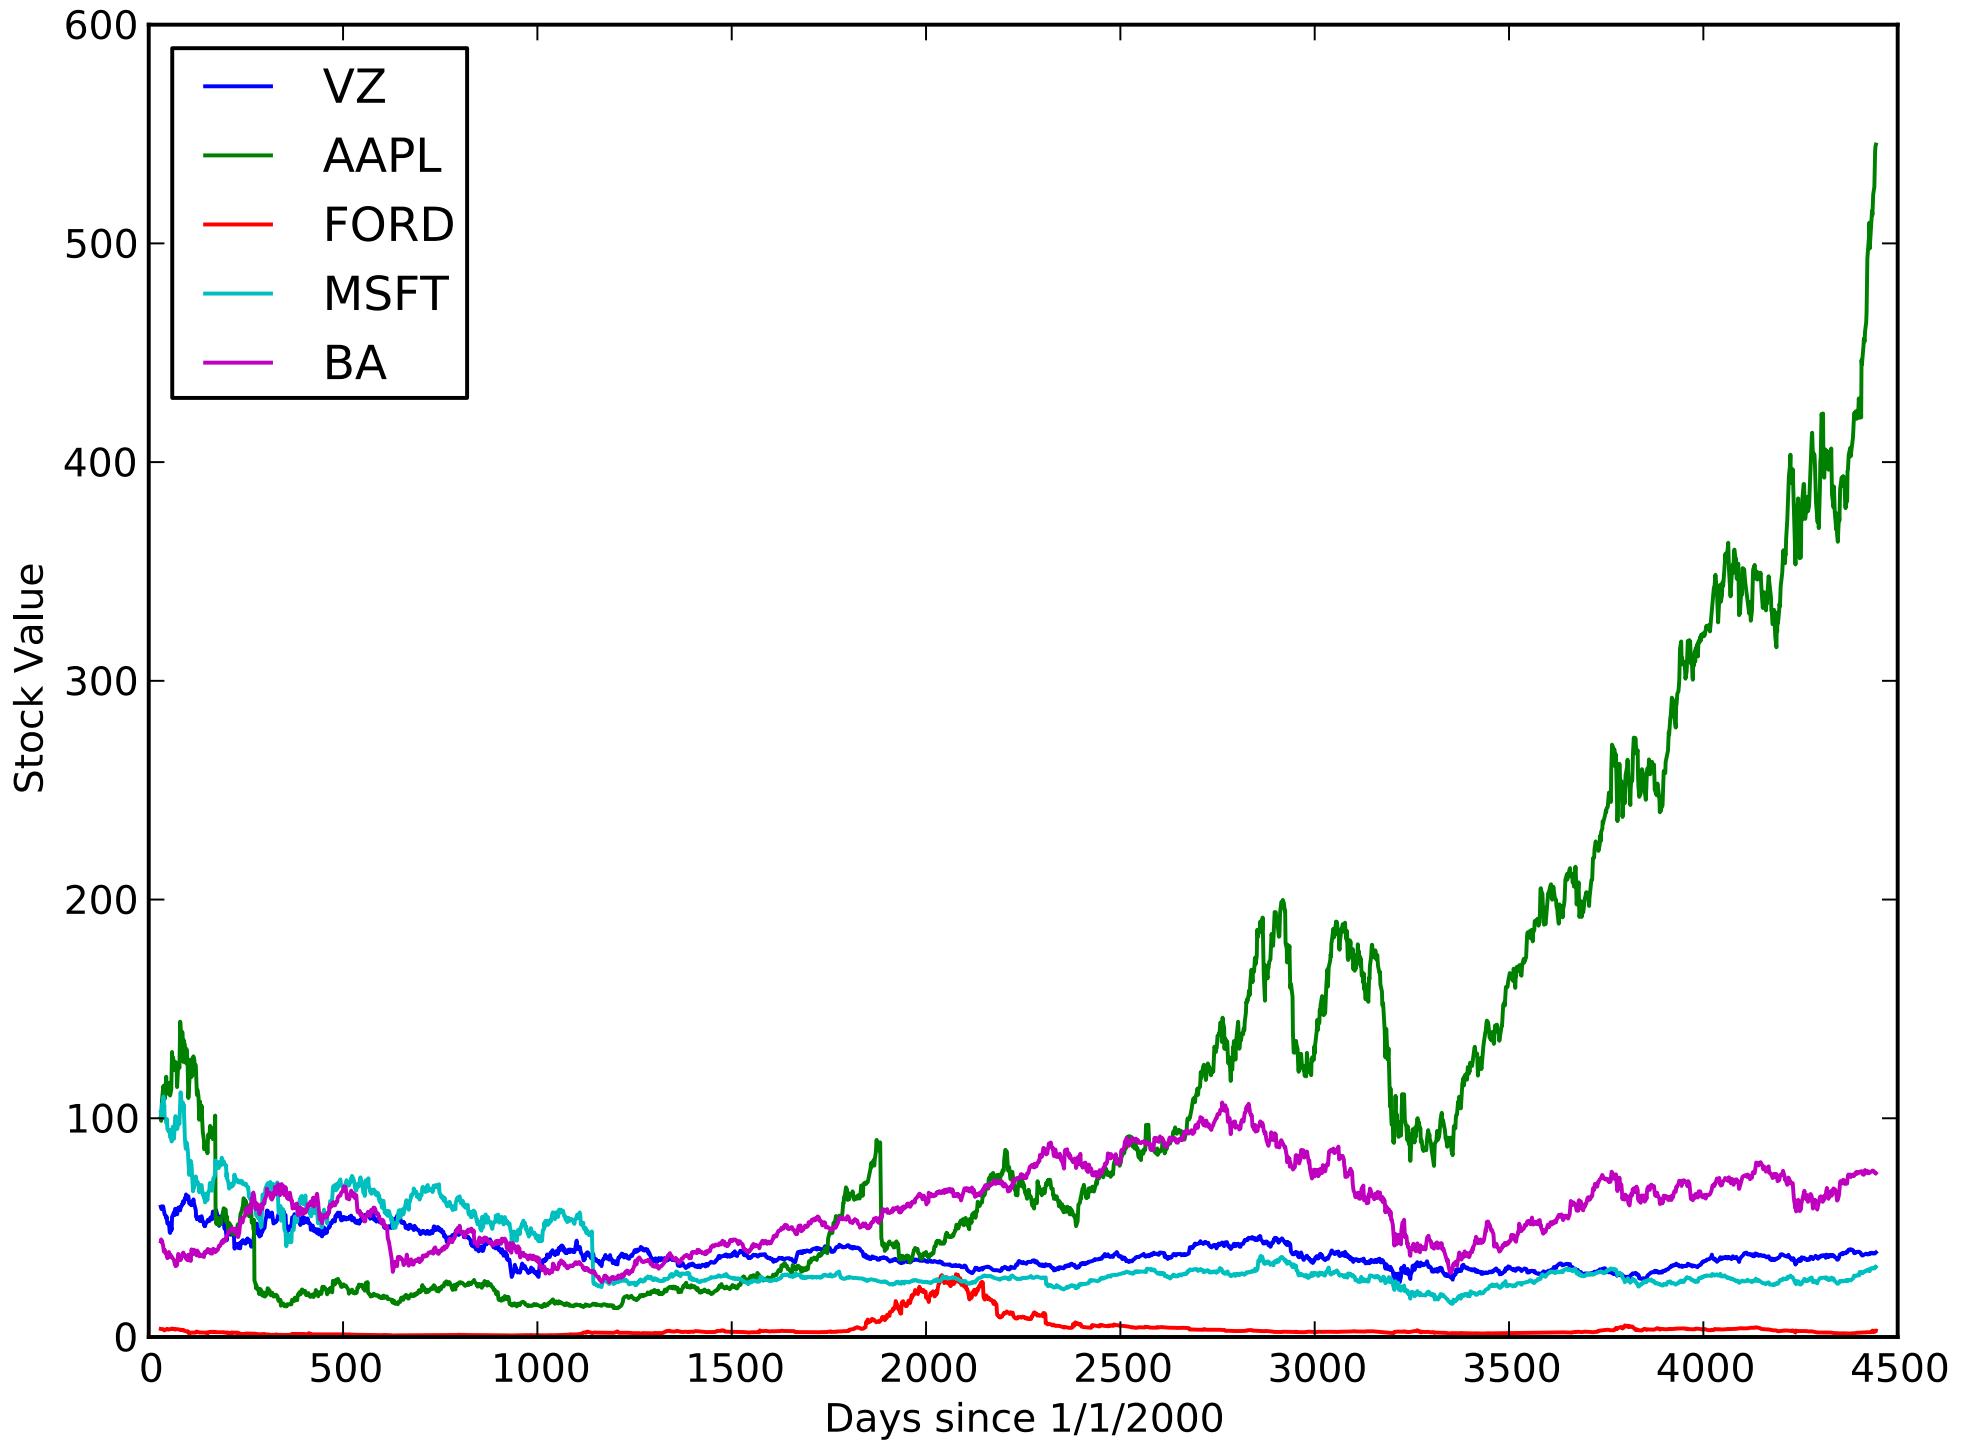
sqlite3

example with plotting data retrieval



To the Notebook!

Stock Performance



MySQLdb

“**MySQLdb** is a thread-compatible interface to the popular MySQL database server that provides the Python database API.”

Not built-in; get it from:

<http://sourceforge.net/projects/mysql-python/>

Project page:

<http://mysql-python.sourceforge.net/MySQLdb.html>

Need MySQL installed, can be tricky on some distributions.

Common syntax with **sqlite3**.

```
sudo brew install mysql; pip install mysql-python
```

MySQLdb

exposing the database

```
import MySQLdb
conn = MySQLdb.connect(host="host", user="user",
    passwd="password", db="database",
    connect_timeout=60)
cursor = conn.cursor()
query_command = "SELECT * FROM table"
cursor.execute(query_command)
db_info = cursor.fetchall()
cursor.close()
conn.close()
```

MySQLdb translates the data-types

Above example is a query, but any MySQL
command can be sent
(respecting user's permissions)

PyGreSQL and Others

Analogous to [MySQLdb](#) for PostgreSQL.

[PyGreSQL](#) - Unix/Windows, Python 2.3-2.6

[pg8000](#) - any OS, Python 2.5+ & 3.0+

[py-postgresql](#) - any OS, Python 3.0+

Breakout Exercise

- 1) Plot Mercator world map using latitude and longitude points from “world.txt”. Polygons are separated by rows containing “nan nan”.
- 2) Import data from “stations_list.txt” into a local database. From now on, retrieve data from the local database before plotting.
- 3) Plot the stations on the Mercator map. Use different colors for Open, Reserved, and Closed. Explore plotting different station networks with different colors.
Extra Credit) Colorize Open stations as a function of elevation.

Breakout Solution

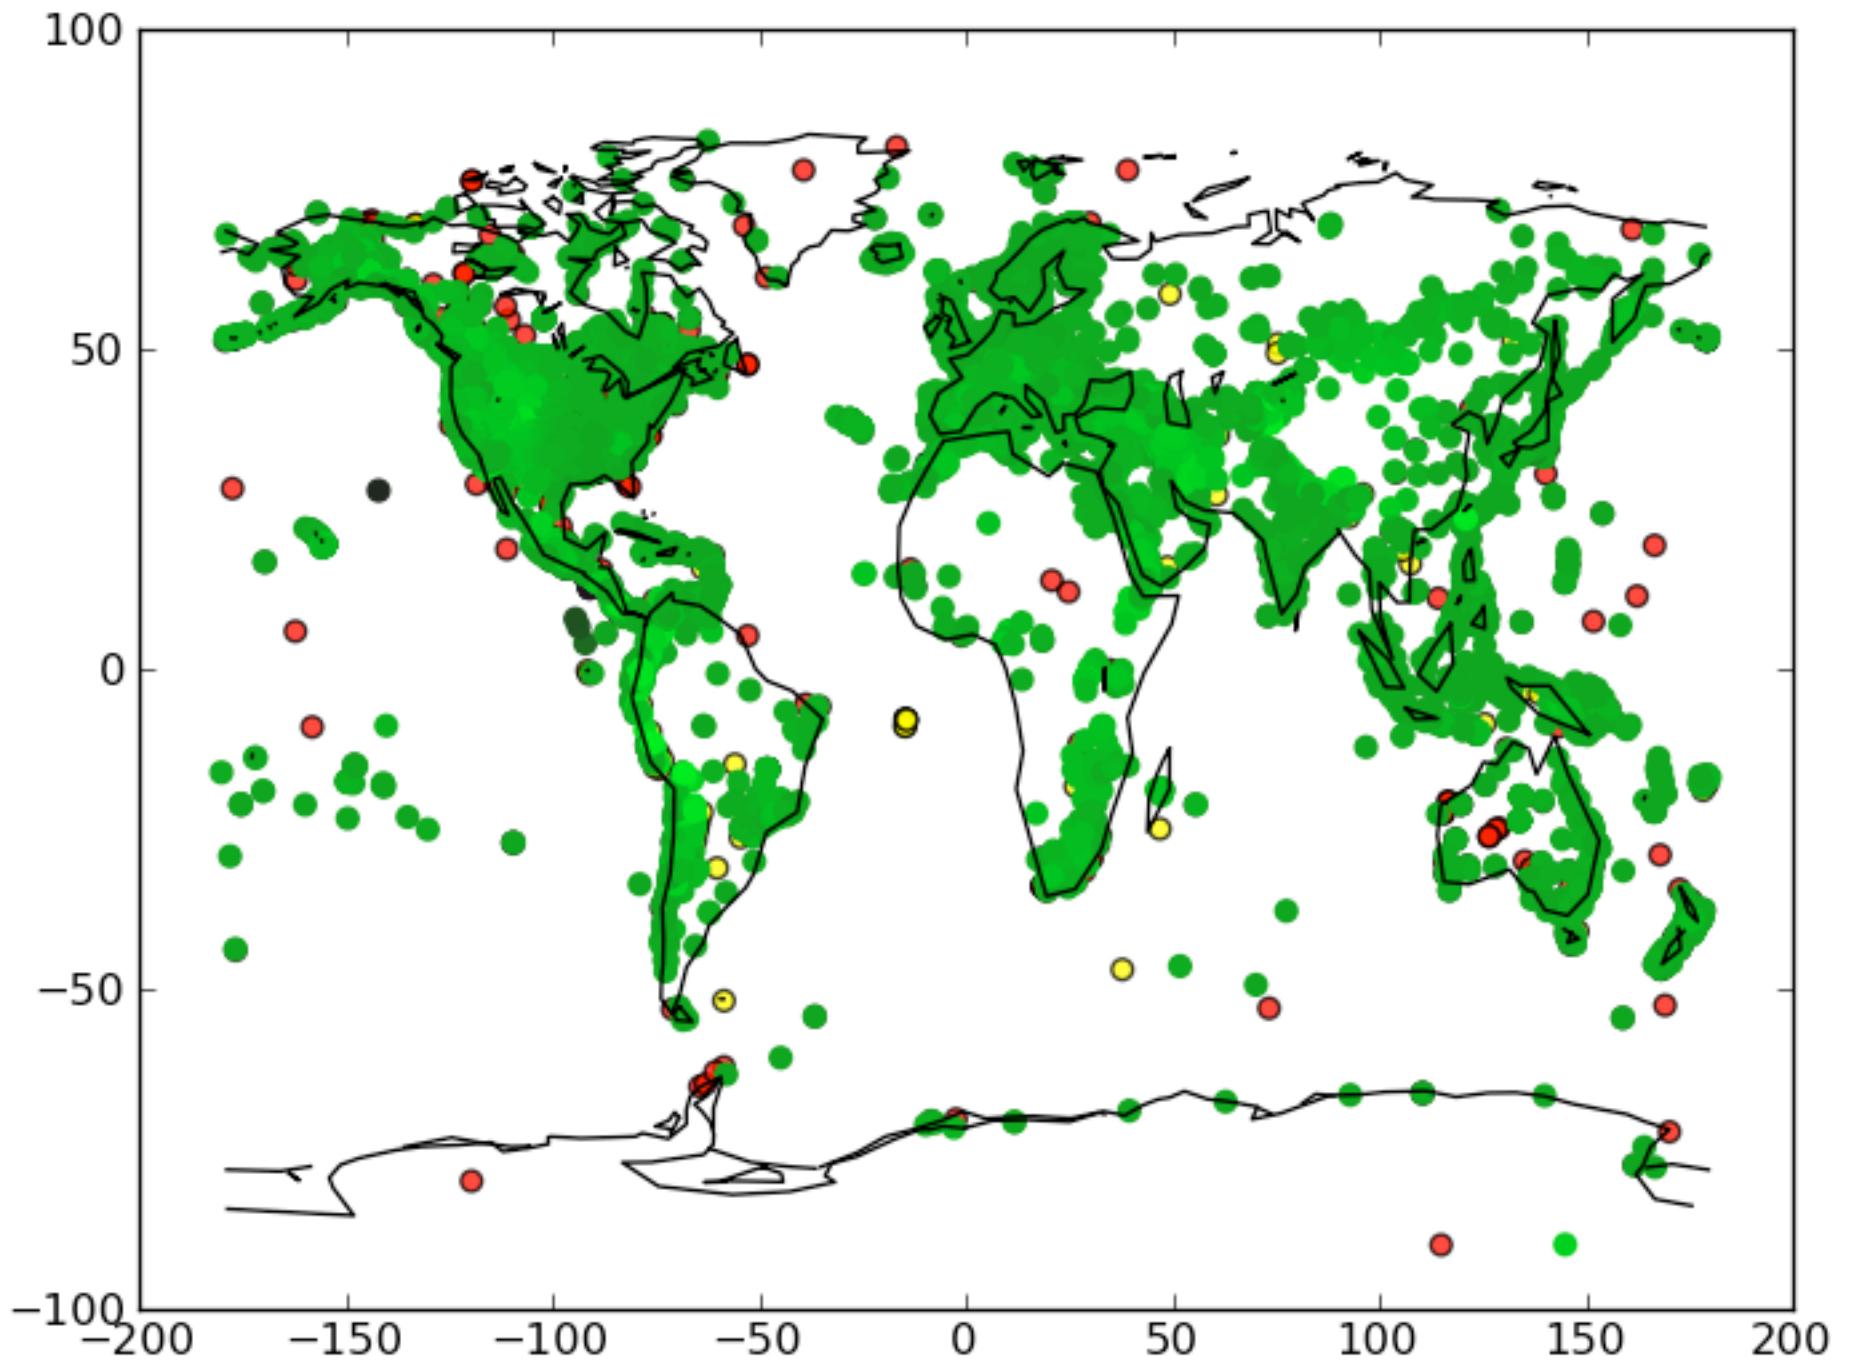
Using a database allows you to avoid the work of separating out the data in python before plotting in different colors or styles. Naively, this could be done with a loop and logic test

```
for station in station_list:  
    if "closed":  
        color="red"
```

but, we can be smarter with a database.



To the Notebook!



HI, THIS IS
YOUR SON'S SCHOOL.
WE'RE HAVING SOME
COMPUTER TROUBLE.



OH, DEAR - DID HE
BREAK SOMETHING?

IN A WAY -)

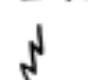


DID YOU REALLY
NAME YOUR SON
Robert'); DROP
TABLE Students;-- ?



OH, YES. LITTLE
BOBBY TABLES,
WE CALL HIM.

WELL, WE'VE LOST THIS
YEAR'S STUDENT RECORDS.
I HOPE YOU'RE HAPPY.



AND I HOPE
YOU'VE LEARNED
TO SANITIZE YOUR
DATABASE INPUTS.

SQLAlchemy

“Object Relational Mapper”

Drizzle | Firebird | Informix | MaxDB | Microsoft Access | Microsoft SQL Server | MySQL | Oracle | PostgreSQL | SQLite | Sybase

- associates Python classes with database tables, and instances of those classes (objects) with rows in their corresponding tables.
- “high level and abstracted pattern of usage”
- essentially, making the Python usage almost entirely independent of the database engine

SQLAlchemy

<http://docs.sqlalchemy.org/en/latest/index.html>

<http://techspot.zenzeek.org/>

Supports

- SQLite
- Postgresql
- MySQL
- Oracle
- MS-SQL
- Firebird
- Sybase
- ...

To the Notebook!





National Energy Research Scientific Computing Center

A DOE Office of Science User Facility
at Lawrence Berkeley National Laboratory

[Site Map](#) | [Help](#) | [Search](#)

Go
[Login](#)

[Home](#)

[About](#)

[News & Media](#)

[Systems](#)

[Support & Services](#)

[Science & Tech](#)

HPC USERS

[MOTD](#)

[Getting Help](#)

Analytics

[Data Analysis & Mining](#)

[Data Exploration](#)

[Data Management](#)

[Visualization](#)

[Workflow Management](#)

[DaVinci: Analytics Server](#)

Systems

[Franklin](#)

[Bassi](#)

[Jacquard](#)

[DaVinci](#)

[PDSF](#)

[HPSS](#)

[Global File System](#)

[Network](#)

[Servers](#)

Services

[Accounts & Allocations](#)

[Announcements](#)

[Consulting](#)

[Grid Computing](#)

[Security](#)

Analytics: Scientific Data Management

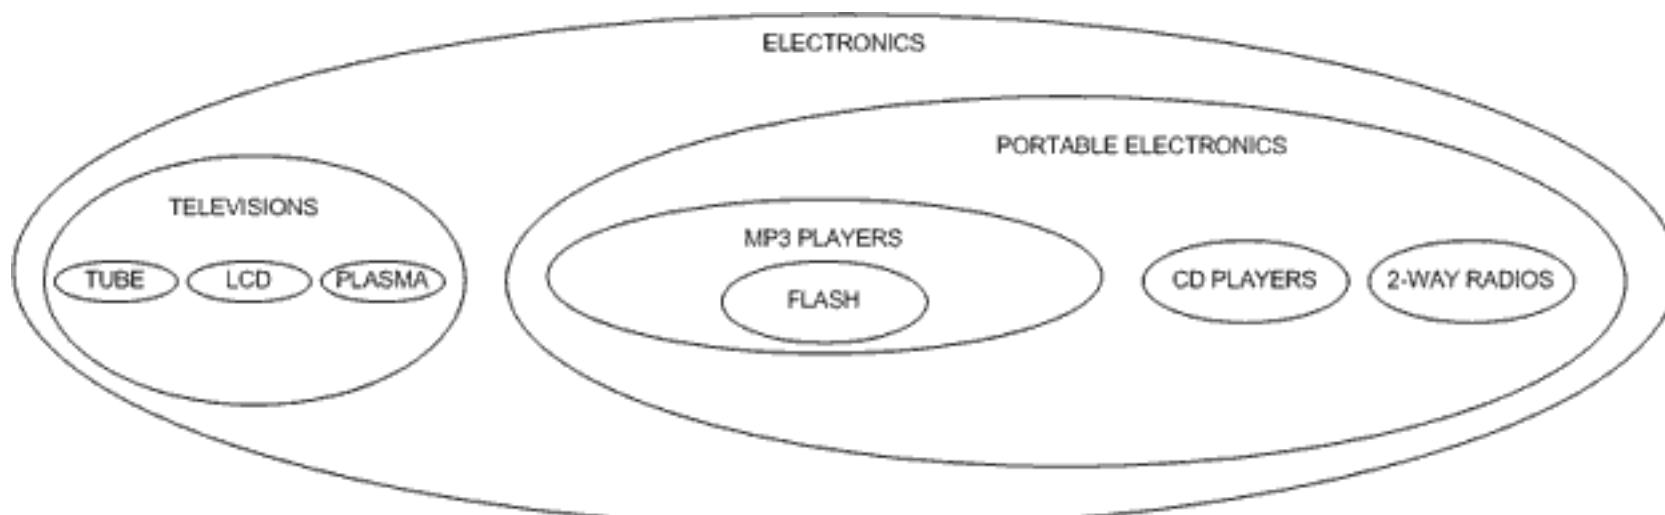
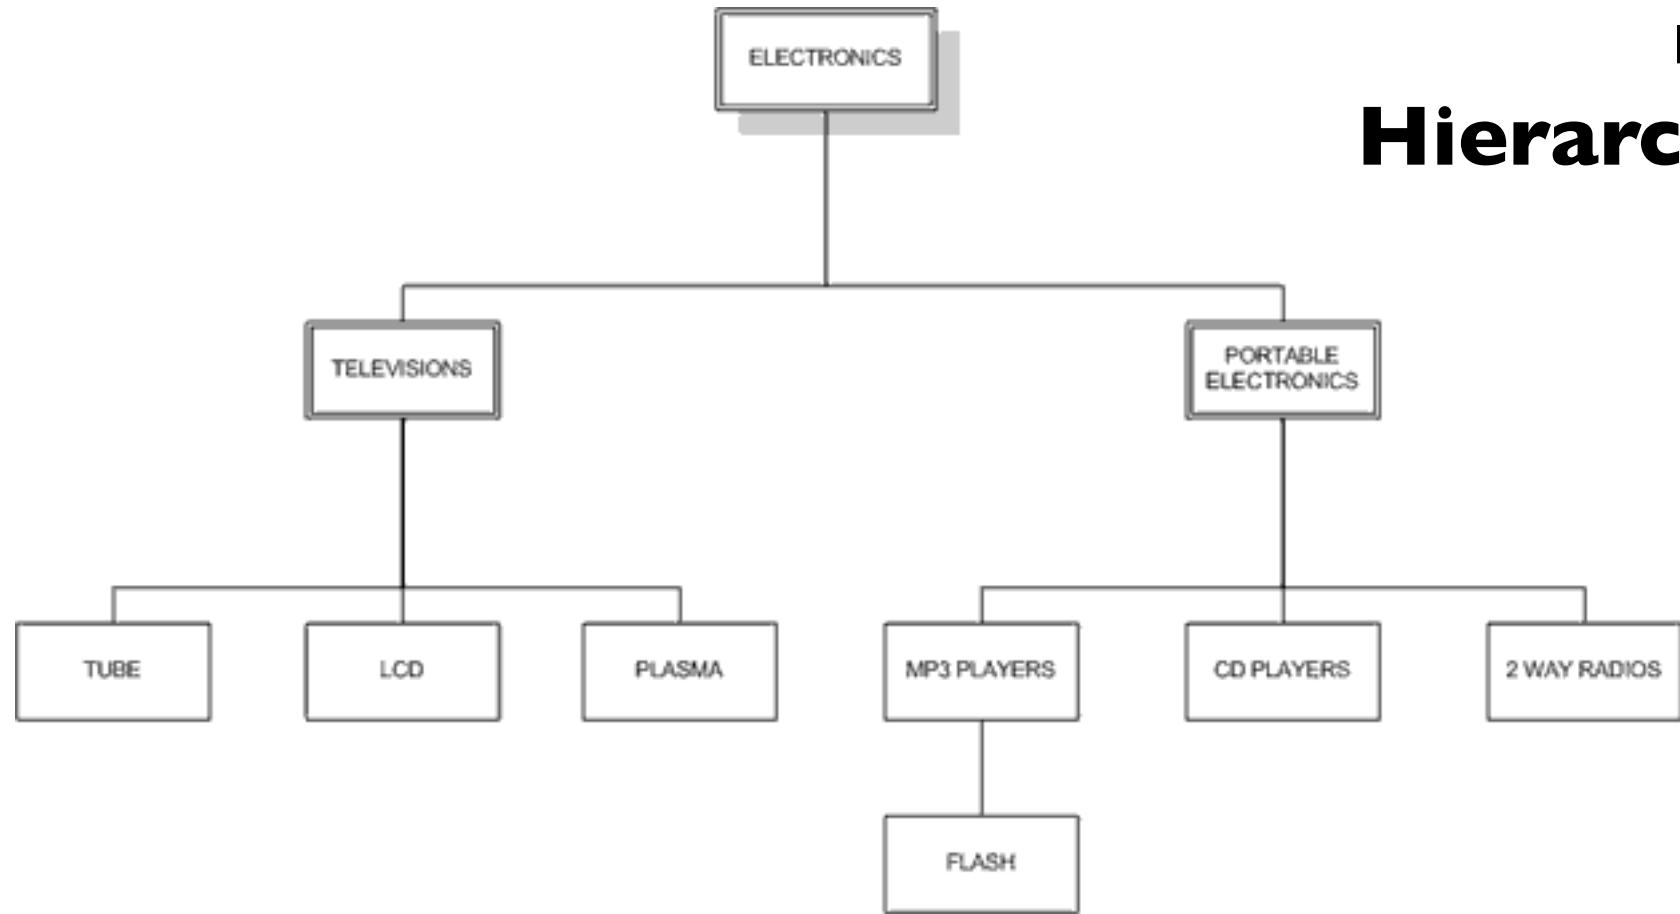
Scientific Data Formats

Scientific data is often based on numerical arrays that are stored in scientific data formats such as [NetCDF](#) (network common data format), [HDF5](#) (hierarchical data format) and [FITS](#) (Flexible Image Transport System). The data formats support efficient mechanisms for accessing and manipulating arrays and allow machine-independent data exchange. They also provide mechanisms for storing metadata information such as simulation parameters.

HDF5 supports parallel I/O which allows multiple processors to read and write files. This is particularly important for large-scale data intensive applications. NetCDF, on the other hand, only supports serial I/O. Parallel I/O is provided by [ParallelNetCDF](#).

Currently there is no common scientific data format that is used across different scientific disciplines. For instance, NetCDF commonly is used in the climate modeling community. FITS is the preferred format for storing astronomy data, while HDF5 is used, for instance, in the combustion and fusion simulation communities.

recall:
Hierarchical



C-based API "database," developed for scientific applications by NCSA

Views data like a file system (with paths /food/spam/ingredients). You connect to the data on disk (like in sqlite3)

Two types of objects in HDF5:

- Datasets - multidimensional homogeneous arrays. They can have attributes (metadata).
 - Table: like record-arrays in numpy
 - Array
- Groups - containers holding datasets and other groups (like "folders")

Several Nice Things about HDF5 itself:

- portability - document is self-describing and platform independent (endian issues are dealt with client-side)
- hugely scalable datamodel (2^{**63} -sized tables)
- Clients made to run massively parallel

pytables is a Pythonic API of HDF5

- happy to use numpy arrays
- works on compressed datafiles
- open-sourced

/

attr

tables.openFile
/ is the root

/Datasets attr

attr

tables.createGroup

/Datasets attr
myarray attr

attr

tables.createArray
*note: different flavors of arrays
unlike numpy arrays, HDF5 arrays can be extended*

/Datasets attr
myarray attr
mytable attr

attr

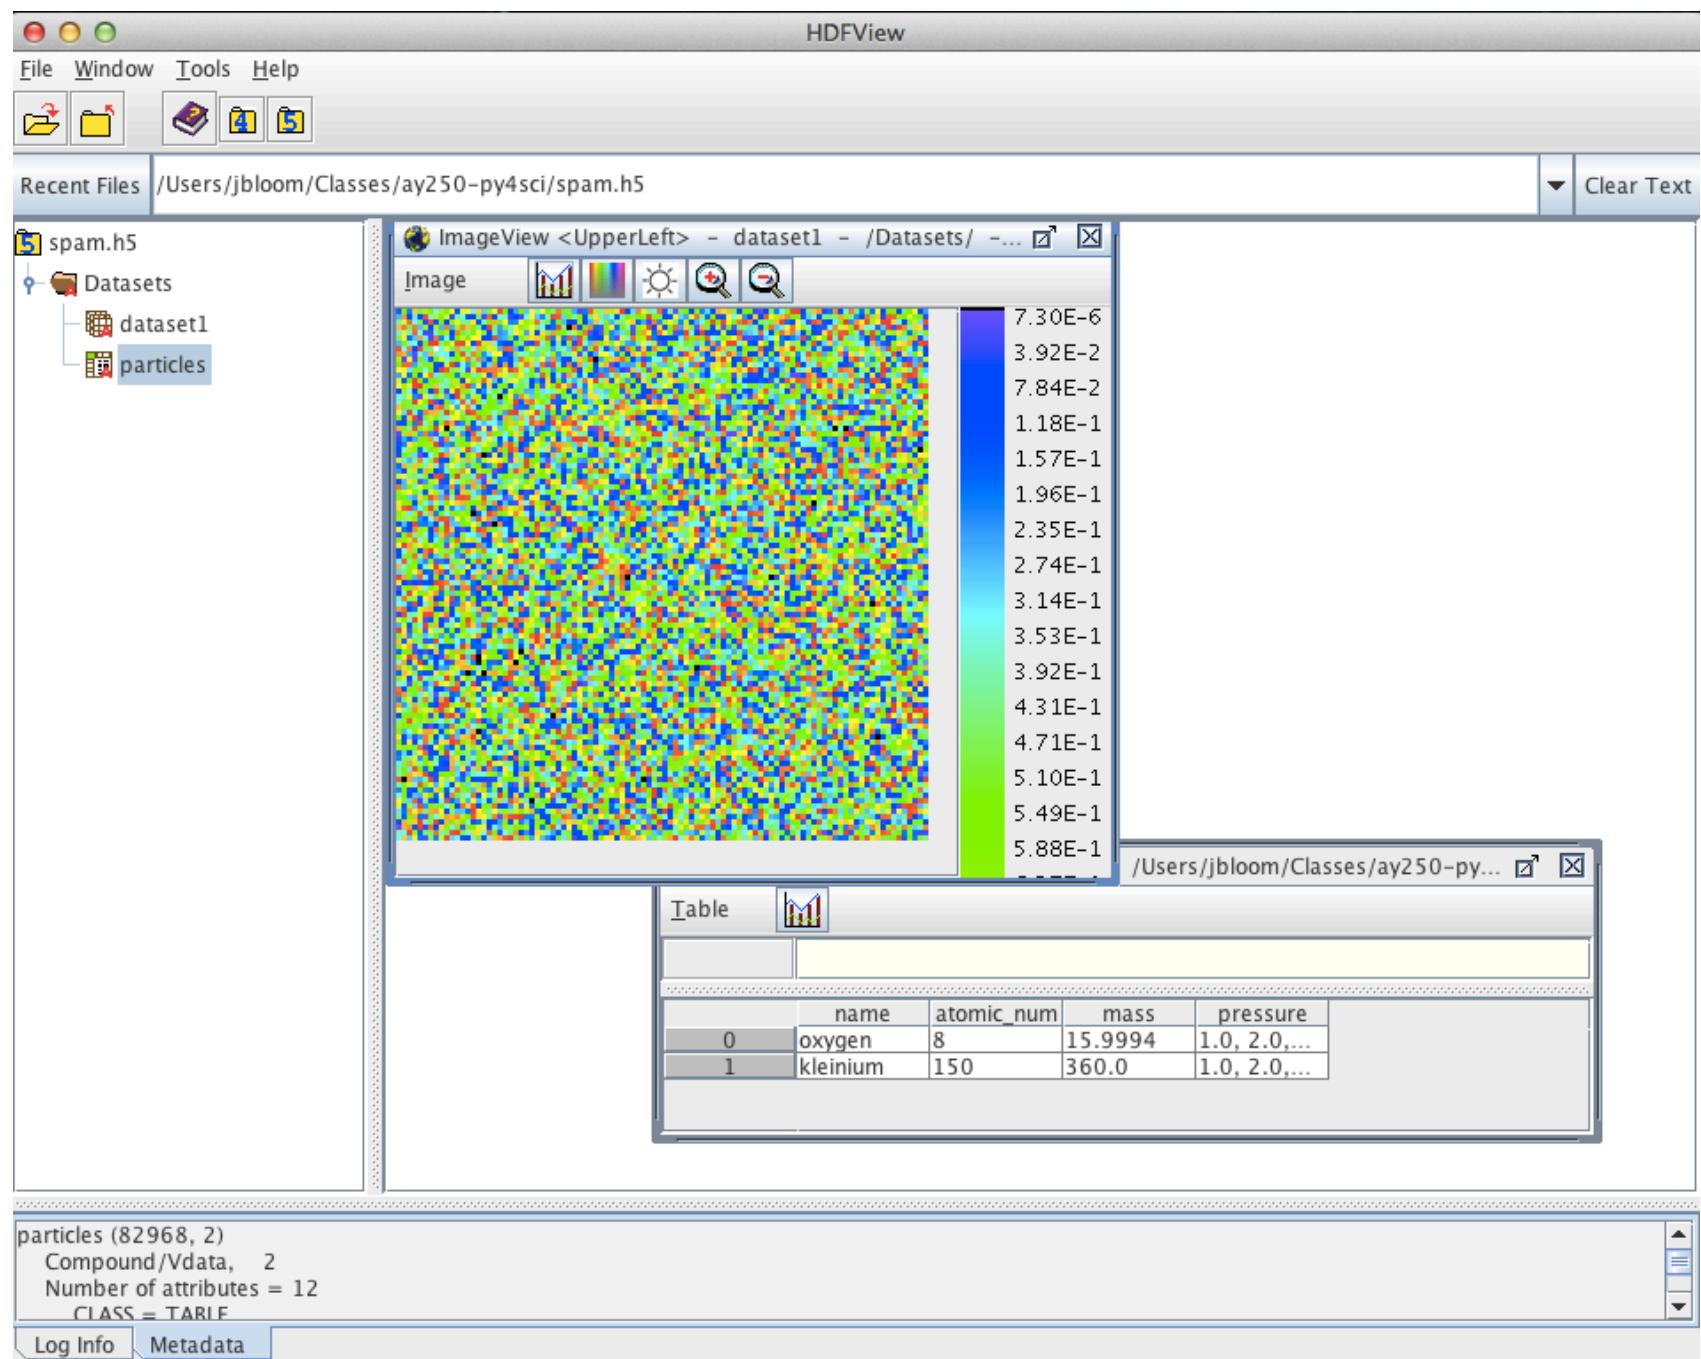
tables.createTable

Also `createExternalLink` creates a link to another HDF5

Making an HDF5 file

```
>>> from tables import *
>>> h5file = openFile("spam.h5", mode = "w", title = "PyTables/HDF5 test file")
>>> h5file
>>> datasets = h5file.createGroup(h5file.root, "Datasets", "Test data sets")
>>> h5file.createArray(datasets, 'dataset1', np.random.random((100,100)), "Test data set #1")
>>> class Particle(IsDescription):
    name      = StringCol(16, pos=1) # 16-character String
    atomic_num = IntCol(pos=2)       # integer
    mass      = FloatCol(pos=3)      # double (double-precision)
    pressure   = Float32Col(shape=(2,3))
>>> table1 = h5file.createTable(datasets, "particles", Particle)
>>> row = table1.row
>>> row
/Datasets/particles.row (Row), pointing to row #0
>>> row["name"] = "oxygen"
>>> row["atomic_num"] = 8
>>> row["mass"] = 15.9994
>>> row["pressure"] = [[1,2,3],[-1,1,3]]
>>> row.append() ; table1.flush()
>>> h5file.root.Datasets.particles[0]
('oxygen', 8, 15.9994, [[1.0, 2.0, 3.0], [-1.0, 1.0, 3.0]])
>>> [row['name'] for row in table1.where('atomic_num > 5 and mass > 5.0')]
['oxygen']

>>> h5file.close()
```



<http://www.hdfgroup.org/hdf-java-html/hdfview/>



“Python and HDF5 - Fast Storage for Large Data”

<https://us.pycon.org/2012/schedule/presentation/231/>

netCDF4

another hierarchical dataformat
(used primarily in climate modeling)

OPeNDAP

Open-source Project for a Network Data Access Protocol

protocols for interacting with remote sites serving
netCDF4, HDF5, etc. data.

Python's netCDF4 module allows you to open
URLs and interact with the data as if it was local

Fast, Distributed DBs (e.g., HBase)



Search wiki, mailing lists & more

Search

HBase Project

Overview

License
Downloads
Release Notes
Issue Tracking
Mailing Lists
Source Repository
Team
Sponsors

Documentation

Getting Started
API
X-Ref
Ref Guide (multi-page)
Ref Guide (single-page)
FAQ
Videos/Presentations
Wiki
ACID Semantics
Bulk Loads
Metrics
HBase on Windows
Cluster replication
Pseudo-Dist. Extras

Welcome to Apache HBase!

HBase is the [Hadoop](#) database. Think of it as a distributed scalable Big Data store.

When Would I Use HBase?

Use HBase when you need random, realtime read/write access to your Big Data. This project's goal is the hosting of very large tables -- billions of rows X millions of columns -- atop clusters of commodity hardware. HBase is an open-source, distributed, versioned, column-oriented store modeled after Google's [Bigtable: A Distributed Storage System for Structured Data](#) by Chang et al. Just as Bigtable leverages the distributed data storage provided by the Google File System, HBase provides Bigtable-like capabilities on top of Hadoop and HDFS.

Features

HBase provides:

- Linear and modular scalability.
- Strictly consistent reads and writes.
- Automatic and configurable sharding of tables
- Automatic failover support between RegionServers.
- Convenient base classes for backing Hadoop MapReduce jobs with HBase tables.
- Easy to use Java API for client access.
- Block cache and Bloom Filters for real-time queries.
- Query predicate push down via server side Filters
- Thrift gateway and a REST-ful Web service that supports XML, Protobuf, and binary data encoding options
- Extensible jruby-based (JIRB) shell
- Support for exporting metrics via the Hadoop metrics subsystem to files or Ganglia; or via JMX

Column-store DBs

large data amenable to map-reduce problems

The screenshot shows the official website for MonetDB at www.monetdb.org/Home. The page features a large blue header with the word "monetdb" and a navigation bar with links for Home, Download, Documentation, Developers, and About us. A sidebar on the left contains links for MonetDB solutions, Getting started, If you need more, and Ongoing development. The main content area discusses the column-store pioneer, highlights column-store features, download and documentation, and ongoing development.

www.monetdb.org/Home

monetdb

Home Download Documentation Developers About us

Overview

MonetDB solutions

- [Column-store features](#)
- [Historic background](#)
- [Science library](#)
- [Project gallery](#)

Getting started

- [Download area](#)
- [Quick start guide](#)

If you need more

- [User guide](#)
- [Reference manuals](#)
- [Common tasks](#)
- [Professional services](#)

Ongoing development

- [Projects and vacancies](#)
- [Nightly testing](#)
- [Bug tracker](#)
- [HG web](#)

The column-store pioneer

The preceding years have been an exciting period. Since 2011 column store technology as pioneered in MonetDB has found its way into the product offerings of all major commercial database vendors. The market for applications empowered by these techniques provide ample space for further innovation, e.g. as demonstrated by [our ongoing projects](#). At the same time, the landscape for major innovations remain wide open. A peek preview is given in the award winning VLDB 2011 paper titled: [The Researcher's Guide to the Data Deluge: Querying a Scientific Database in Just a Few Seconds](#).

Column-store features

When your database grows into millions of records spread over lots of tables and used in business or science data warehouse applications, you really want a column-store database management system.

MonetDB innovates at all layers of a DBMS, e.g. a storage model based on vertical fragmentation, a modern CPU-tuned query execution architecture, automatic and self-

Download and documentation

The download section contains information to get started with the binary or source download. Once installed, you may want to follow the short tutorials for Linux or Windows, or explore the SQL documentation for details about the language and recipes for common tasks in application development.

The nuts and bolts on the database kernel are described in the

Ongoing development

MonetDB is actively used in our research and real life applications. Nightly builds and regression testing ensure its quality, bug tracking helps to collect experiences and feature requests. Browsing the source code repository is supported by the Mercurial web frontend. Contributions ranging from bug reports, cross-platform issues, patches and features are highly



Get Latest SciDB



SHARE

ABOUT | USE CASES | GET INVOLVED | DOWNLOAD | NEWS & EVENTS | FORUM

FOR THE TOUGHEST PROBLEMS ON THE PLANET

SciDB Open Source Data Management and
Analytics Software for Scientific Research

HEART BEAT: Oct 23, 2011: Slides from the 2nd SciDB Community Meeting have been posted

June 15, 2011: Release 11.06 is Available

Copyright 2010-2011, SciDB.org

Website Design by MicroArts Creative Agency

HI, THIS IS
YOUR SON'S SCHOOL.
WE'RE HAVING SOME
COMPUTER TROUBLE.



OH, DEAR - DID HE
BREAK SOMETHING?

IN A WAY -)



DID YOU REALLY
NAME YOUR SON
Robert'); DROP
TABLE Students;-- ?



OH, YES. LITTLE
BOBBY TABLES,
WE CALL HIM.

WELL, WE'VE LOST THIS
YEAR'S STUDENT RECORDS.
I HOPE YOU'RE HAPPY.



AND I HOPE
YOU'VE LEARNED
TO SANITIZE YOUR
DATABASE INPUTS.